



МИНОБРНАУКИ РОССИИ  
Федеральное государственное бюджетное образовательное учреждение  
высшего образования  
**«МИРЭА – Российский технологический университет»**  
**РТУ МИРЭА**

---

ИКБ направление «Киберразведка и противодействие угрозам с применением технологий искусственного интеллекта» 10.04.01

Кафедра КБ-4 «Интеллектуальные системы информационной безопасности»

**Лабораторная работа №2**

по дисциплине

«Анализ защищённости систем искусственного интеллекта»

Группа:  
ББМО-01-22  
Выполнила:  
Огольцова Н.Д.

Проверил:  
Спирин А.А.

Москва 2023

## Ход работы: (Раздел 1)

### 1) Установим необходимые библиотеки и фреймворки;

```
!pip install adversarial-robustness-toolbox
```

```
Requirement already satisfied: adversarial-robustness-toolbox in /usr/local/lib/python3.10/dist-packages (1.16.0)
Requirement already satisfied: numpy>=1.18.0 in /usr/local/lib/python3.10/dist-packages (from adversarial-robustness-toolbox) (1.23.5)
Requirement already satisfied: scipy>=1.4.1 in /usr/local/lib/python3.10/dist-packages (from adversarial-robustness-toolbox) (1.11.3)
Requirement already satisfied: scikit-learn<1.2.0,>=0.22.2 in /usr/local/lib/python3.10/dist-packages (from adversarial-robustness-toolbox) (1.1.3)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from adversarial-robustness-toolbox) (1.16.0)
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from adversarial-robustness-toolbox) (67.7.2)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from adversarial-robustness-toolbox) (4.66.1)
Requirement already satisfied: joblib>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn<1.2.0,>=0.22.2->adversarial-robustness-toolbox) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn<1.2.0,>=0.22.2->adversarial-robustness-toolbox) (3.2.0)
```

```
import cv2
import os
import torch
import random
import pickle
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from sklearn.model_selection import train_test_split
from keras.utils import to_categorical
from keras.applications import ResNet50
from keras.applications import VGG16
from keras.applications.resnet50 import preprocess_input
from keras.preprocessing import image
from keras.models import load_model, save_model
from keras.layers import Dense, Flatten, GlobalAveragePooling2D
from keras.models import Model
from keras.optimizers import Adam
from keras.losses import categorical_crossentropy
from keras.metrics import categorical_accuracy
from keras.callbacks import ModelCheckpoint, EarlyStopping, TensorBoard
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPool2D, AvgPool2D, BatchNormalization, Reshape, Lambda
from art.estimators.classification import KerasClassifier
from art.attacks.evasion import FastGradientMethod, ProjectedGradientDescent

import keras
import tensorflow as tf
from keras.applications.vgg16 import VGG16
from keras.applications.resnet50 import ResNet50
from keras.applications.mobilenet import MobileNet
from keras.models import Model
from keras.preprocessing import image
from tensorflow.keras.layers import Input, Lambda, Dense, Flatten, Dropout
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import os
import cv2
import random
from tensorflow.keras.preprocessing.image import ImageDataGenerator, img_to_array, array_to_img, load_img

from tensorflow.keras.utils import to_categorical
from tensorflow import keras as k

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import datasets, transforms

import zipfile
from google.colab import drive
from PIL import Image
from skimage import io
import pandas as pd
import matplotlib.image as img
from PIL import Image
from sklearn import svm
from sklearn.model_selection import train_test_split
```

### 2) Подключим Google Drive для работы;

```

ls
drive.mount('/content/drive/')

drive  Meta      ResNet50.h5  test  Test.csv  Train
meta   Meta.csv  sample_data  Test  train     Train.csv
Drive already mounted at /content/drive/; to attempt to forcibly remount, call drive.mount("/content/drive/", force_remount=True).

```

### 3) Разархивируем файл с датасетом;

```

zip_file = '/content/drive/MyDrive/A3CMI/archive.zip'
z = zipfile.ZipFile(zip_file, 'r')
z.extractall()

print(os.listdir())

['.config', 'test', 'meta', 'Test.csv', 'Train.csv', 'Train', 'Meta.csv', 'drive', 'Meta', 'Test', 'train', 'ResNet50.h5', 'sample_data']

```

### 4) Зададим переменные для удобной навигации по датасету;

```

data_path = '/content'
train_data_path = os.path.join(data_path, 'Train')
test_data_path = os.path.join(data_path, 'Test')
meta_data_path = os.path.join(data_path, 'Meta')

```

### 5) С помощью цикла преобразуем и сохраним изображения в виде массива в переменную;

```

data = []
labels = []
CLASSES = 43
# Используем цикл for для доступа к каждому изображению
for i in range(CLASSES):
    img_path = os.path.join(train_data_path, str(i)) #0-42
    for img in os.listdir(img_path):
        im = Image.open(train_data_path + '/' + str(i) + '/' + img)
        im = im.resize((32,32))
        im = np.array(im)
        data.append(im)
        labels.append(i)
data = np.array(data)
labels = np.array(labels)
print("data[0]: ", data[0])
print("labels[0: ]", labels[0])

data[0]: [[255 255 255]
 [255 255 255]
 [255 255 255]
 ...
 [255 255 255]
 [255 255 255]
 [255 255 255]]

[[255 255 255]

```

### 6) Разделим датасет на обучающую и тестовую выборки;



Скомпилируем её:

```
model.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["accuracy"])
history = model.fit(x_train, y_train, epochs=5, batch_size=64, validation_data=(x_test, y_test))

Train on 31367 samples, validate on 7842 samples
Epoch 1/5
31367/31367 [=====] - ETA: 0s - loss: 1.6415 - accuracy: 0.5378/usr/local/lib/python3.10/dist-packages/keras/src/er
updates = self.state_updates
31367/31367 [=====] - 75s 2ms/sample - loss: 1.6415 - accuracy: 0.5378 - val_loss: 4.7476 - val_accuracy: 0.1750
Epoch 2/5
31367/31367 [=====] - 62s 2ms/sample - loss: 0.4007 - accuracy: 0.8812 - val_loss: 3.7772 - val_accuracy: 0.2548
Epoch 3/5
31367/31367 [=====] - 71s 2ms/sample - loss: 0.2662 - accuracy: 0.9220 - val_loss: 3.2438 - val_accuracy: 0.3284
Epoch 4/5
31367/31367 [=====] - 63s 2ms/sample - loss: 0.1945 - accuracy: 0.9468 - val_loss: 1.1653 - val_accuracy: 0.7308
Epoch 5/5
31367/31367 [=====] - 65s 2ms/sample - loss: 0.1489 - accuracy: 0.9594 - val_loss: 1.0241 - val_accuracy: 0.7904
```

8)Создадим модель ResNet50;

```
resnet = k.applications.ResNet50(weights='imagenet', include_top=False)

model_2 = k.models.Sequential([
    resnet,
    tf.keras.layers.GlobalAveragePooling2D(),
    k.layers.Dropout(0.2),
    k.layers.Dense(256, activation='relu'),
    k.layers.BatchNormalization(),
    k.layers.Dropout(0.1),
    k.layers.Dense(512, activation='relu'),
    k.layers.BatchNormalization(),
    k.layers.Dropout(0.2),
    k.layers.Dense(43, activation='softmax')
])
print(model_2.summary())
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
=====		
resnet50 (Functional)	(None, None, None, 2048)	23587712
global_average_pooling2d_4 (GlobalAveragePooling2D)	(None, 2048)	0
dropout_12 (Dropout)	(None, 2048)	0
dense_12 (Dense)	(None, 256)	524544
batch_normalization_8 (BatchNormalization)	(None, 256)	1024
dropout_13 (Dropout)	(None, 256)	0
dense_13 (Dense)	(None, 512)	131584
batch_normalization_9 (BatchNormalization)	(None, 512)	2048
dropout_14 (Dropout)	(None, 512)	0
dense_14 (Dense)	(None, 43)	22059
=====		
Total params: 24268971 (92.58 MB)		
Trainable params: 24214315 (92.37 MB)		
Non-trainable params: 54656 (213.50 KB)		
=====		
None		

Скомпилируем её:

```
model_2.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["accuracy"])
history_2 = model_2.fit(x_train, y_train, epochs=5, batch_size=64, validation_data=(x_test, y_test))
```

Train on 31367 samples, validate on 7842 samples

```
Epoch 1/5
31367/31367 [=====] - 97s 3ms/sample - loss: 0.8805 - accuracy: 0.7560 - val_loss: 0.2633 - val_accuracy: 0.9306
Epoch 2/5
31367/31367 [=====] - 67s 2ms/sample - loss: 0.2269 - accuracy: 0.9426 - val_loss: 0.1763 - val_accuracy: 0.9461
Epoch 3/5
31367/31367 [=====] - 70s 2ms/sample - loss: 0.2242 - accuracy: 0.9479 - val_loss: 0.4390 - val_accuracy: 0.9082
Epoch 4/5
31367/31367 [=====] - 71s 2ms/sample - loss: 0.1403 - accuracy: 0.9666 - val_loss: 0.3176 - val_accuracy: 0.9325
Epoch 5/5
31367/31367 [=====] - 68s 2ms/sample - loss: 0.1094 - accuracy: 0.9759 - val_loss: 0.0957 - val_accuracy: 0.9742
```

9)Создадим модель VGG16;

```
vgg16 = k.applications.VGG16(weights='imagenet', include_top=False)

model_3 = k.models.Sequential([
    vgg16,
    tf.keras.layers.GlobalAveragePooling2D(),
    k.layers.Dropout(0.2),
    k.layers.Dense(256, activation='relu'),
    k.layers.BatchNormalization(),
    k.layers.Dropout(0.1),
    k.layers.Dense(512, activation='relu'),
    k.layers.BatchNormalization(),
    k.layers.Dropout(0.2),
    k.layers.Dense(43, activation='softmax')
])
print(model_3.summary())
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
=====	=====	=====
vgg16 (Functional)	(None, None, None, 512)	14714688
global_average_pooling2d_5 (GlobalAveragePooling2D)	(None, 512)	0
dropout_15 (Dropout)	(None, 512)	0
dense_15 (Dense)	(None, 256)	131328
batch_normalization_10 (Ba tchNormalization)	(None, 256)	1024
dropout_16 (Dropout)	(None, 256)	0
dense_16 (Dense)	(None, 512)	131584
batch_normalization_11 (Ba tchNormalization)	(None, 512)	2048
dropout_17 (Dropout)	(None, 512)	0
dense_17 (Dense)	(None, 43)	22059
=====	=====	=====
Total params: 15002731 (57.23 MB)		
Trainable params: 15001195 (57.23 MB)		
Non-trainable params: 1536 (6.00 KB)		

None

Скомпилируем её:

```
model_3.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["accuracy"])
history_3 = model_3.fit(x_train, y_train, epochs=5, batch_size=64, validation_data=(x_test, y_test))

Train on 31367 samples, validate on 7842 samples
Epoch 1/5
31367/31367 [=====] - 25s 799us/sample - loss: 3.2333 - accuracy: 0.1182 - val_loss: 2.5970 - val_accuracy: 0.1854
Epoch 2/5
31367/31367 [=====] - 20s 652us/sample - loss: 2.5899 - accuracy: 0.1820 - val_loss: 2.6285 - val_accuracy: 0.1862
Epoch 3/5
31367/31367 [=====] - 20s 645us/sample - loss: 2.1403 - accuracy: 0.2764 - val_loss: 2.7435 - val_accuracy: 0.1747
Epoch 4/5
31367/31367 [=====] - 20s 646us/sample - loss: 1.4171 - accuracy: 0.4944 - val_loss: 1.0392 - val_accuracy: 0.6551
Epoch 5/5
31367/31367 [=====] - 22s 693us/sample - loss: 0.6955 - accuracy: 0.7597 - val_loss: 0.4497 - val_accuracy: 0.8564
```

10) Полученный результат, представленный в виде таблицы по разделу 1:

Model	Training Accuracy	Validation Accuracy	Test Accuracy
MobileNet	95.9416	79.036	97.5899
Resnet50	97.593	97.4241	97.1946
VGG16	75.9652	85.6414	88.7911

## Задание 2:

1) Для второго задания используем тысячу первых тестовых изображений, также зададим значения epsilon;

```
[13] epsilons = [0.004, 0.008, 0.012, 0.016, 0.02, 0.031, 0.039, 0.078, 0.196, 0.314]
```

```
x_test = data[:1000]
y_test = y_test[:1000]
```


2) Переопределим значения epsilon выведем изображение под номером 14 после атаки FGSM в сравнении с исходным;

```
for eps in epsilons:
    attack_fgsm.set_params(**{'eps': eps})
    x_test_adv = attack_fgsm.generate(x_test, y_test)
    pred = np.argmax(model.predict(x_test_adv[13:14]))


    plt.figure(i)
    plt.title(f"Epsilon: {eps} , Prediction: {pred}, True -> {np.argmax(y_test[0])}")
    plt.imshow(x_test_adv[14])
    plt.show()
    i += 1
```




Epsilon: 0, Prediction: 5, True -> 4



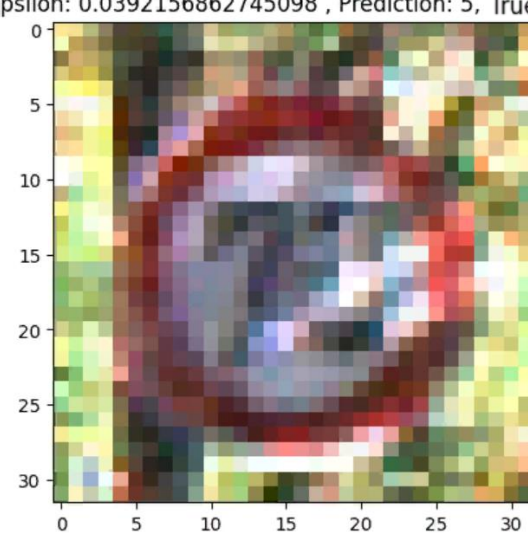
Epsilon: 0.00392156862745098 , Prediction: 5, True -> 4



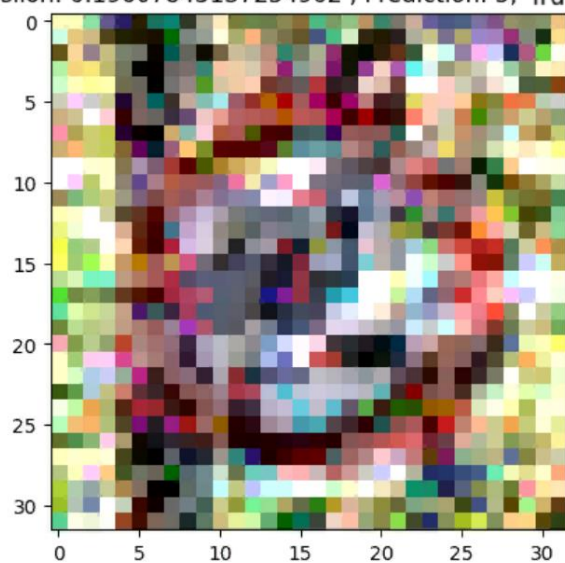
Epsilon: 0.0196078431372549 , Prediction: 5, True -> 4




Epsilon: 0.0392156862745098 , Prediction: 5, True -> 4



Epsilon: 0.19607843137254902 , Prediction: 5, True -> 4




Epsilon: 0.3137254901960784 , Prediction: 5, True -> 4






3) Переопределим значения  $\epsilon$  выведем изображение под номером 16 после атаки FGSM в сравнении с исходным уже для модели VGG16;


Epsilon: 0, Prediction: 15, True -> 16




Epsilon: 0.00392156862745098, Prediction: 16, True -> 16



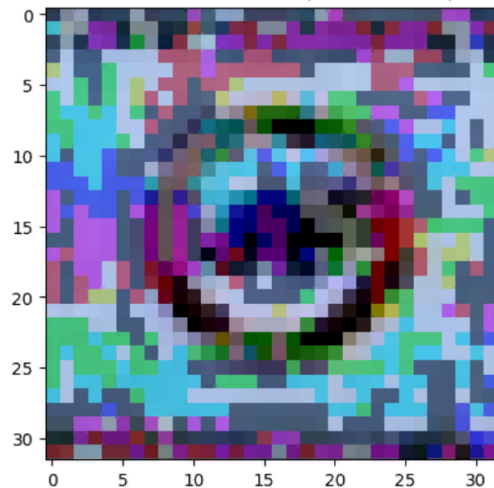
Epsilon: 0.0196078431372549, Prediction: 16, True -> 16



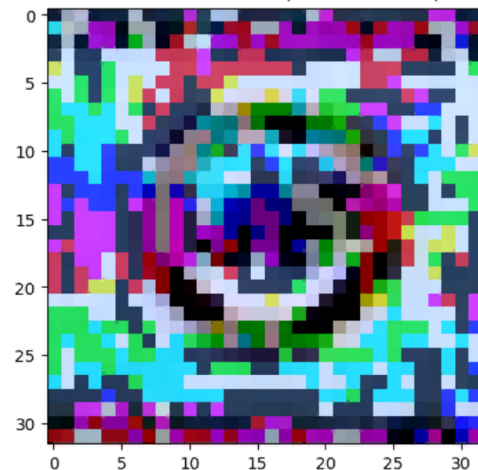
Epsilon: 0.0392156862745098, Prediction: 8, True -> 16




Epsilon: 0.19607843137254902, Prediction: 7, True -> 16




Epsilon: 0.3137254901960784, Prediction: 18, True -> 16



4) Для модели ResNet50 построим график зависимости точности классификации от параметра искажений  $\epsilon$ ;



5) Для модели VGG16 также построим график зависимости точности классификации от параметра искажений  $\epsilon$ ;



6) Сохраним полученные значения точности для каждой из атак, аналогично для обеих моделей;

```

✓ [19] attack_fgsm = FastGradientMethod(estimator=classifier, eps=0)
17 adv accuracies_fgsm = []
ек. x_test_adv = attack_fgsm.generate(x_test, y_test)
loss, accuracy = model.evaluate(x_test_adv, y_test)
adv accuracies_fgsm.append(accuracy)
resnetfgsm_0=accuracy

attack_fgsm = FastGradientMethod(estimator=classifier, eps=1/255)
adv accuracies_fgsm = []
x_test_adv = attack_fgsm.generate(x_test, y_test)
loss, accuracy = model.evaluate(x_test_adv, y_test)
adv accuracies_fgsm.append(accuracy)
resnetfgsm_1=accuracy

attack_fgsm = FastGradientMethod(estimator=classifier, eps=5/255)
adv accuracies_fgsm = []
x_test_adv = attack_fgsm.generate(x_test, y_test)
loss, accuracy = model.evaluate(x_test_adv, y_test)
adv accuracies_fgsm.append(accuracy)
resnetfgsm_2=accuracy

attack_fgsm = FastGradientMethod(estimator=classifier, eps=10/255)
adv accuracies_fgsm = []
x_test_adv = attack_fgsm.generate(x_test, y_test)
loss, accuracy = model.evaluate(x_test_adv, y_test)
adv accuracies_fgsm.append(accuracy)
resnetfgsm_3=accuracy

```

```

attack_pgd = ProjectedGradientDescent(estimator=classifier, eps=0, max_iter=4, verbose=False)
adv accuracies_pgd = []
x_test_adv = attack_pgd.generate(x_test, y_test)
loss, accuracy = model.evaluate(x_test_adv, y_test)
adv accuracies_pgd.append(accuracy)
resnetpgd_0=accuracy

attack_pgd = ProjectedGradientDescent(estimator=classifier, eps=1/255, max_iter=4, verbose=False)
adv accuracies_pgd = []
x_test_adv = attack_pgd.generate(x_test, y_test)
loss, accuracy = model.evaluate(x_test_adv, y_test)
adv accuracies_pgd.append(accuracy)
resnetpgd_1=accuracy

attack_pgd = ProjectedGradientDescent(estimator=classifier, eps=5/255, max_iter=4, verbose=False)
adv accuracies_pgd = []
x_test_adv = attack_pgd.generate(x_test, y_test)
loss, accuracy = model.evaluate(x_test_adv, y_test)
adv accuracies_pgd.append(accuracy)
resnetpgd_2=accuracy

attack_pgd = ProjectedGradientDescent(estimator=classifier, eps=10/255, max_iter=4, verbose=False)
adv accuracies_pgd = []
x_test_adv = attack_pgd.generate(x_test, y_test)
loss, accuracy = model.evaluate(x_test_adv, y_test)
adv accuracies_pgd.append(accuracy)
resnetpgd_3=accuracy

```

7) Выведем получившиеся значения в виде таблицы;

✓

0

ПК

▶

from tabulate import tabulate

```

table = [
    ["Model", "Source Image", "Adversarial images  $\epsilon=1/255$ ", "Adversarial images  $\epsilon=5/255$ ", "Adversarial images  $\epsilon=10/255$ "],
    ["Resnet50 - FGSM", resnetfgsm_0, resnetfgsm_1, resnetfgsm_2, resnetfgsm_3],
    ["Resnet50 - PGD", resnetpgd_0, resnetpgd_1, resnetpgd_2, resnetpgd_3],
    ["VGG16 - FGSM", vggfgsm_0, vggfgsm_1, vggfgsm_2, vggfgsm_3],
    ["VGG16 - PGD", vggpgd_0, vggpgd_1, vggpgd_2, vggpgd_3]]

table1 = tabulate(table, headers="firstrow", tablefmt="grid")
print(table1)

```

Model	Source Image	Adversarial images $\epsilon=1/255$	Adversarial images $\epsilon=5/255$	Adversarial images $\epsilon=10/255$
Resnet50 - FGSM	0.052	0.048	0.042	0.022
Resnet50 - PGD	0.052	0.048	0.043	0.02
VGG16 - FGSM	0.03	0.03	0.004	0.004
VGG16 - PGD	0.03	0.03	0.006	0.002

### Задание 3

1)Загрузим наше изображение знака «Стоп» и будем стараться его определять, как изображение «Ограничение скорости 30»

✓







0

ПК



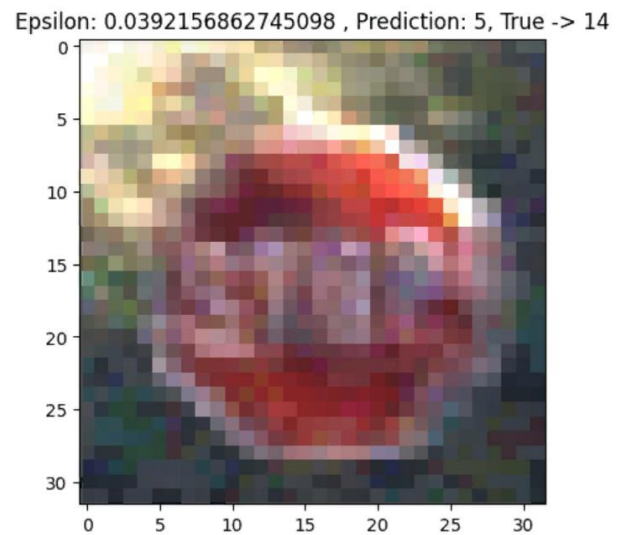
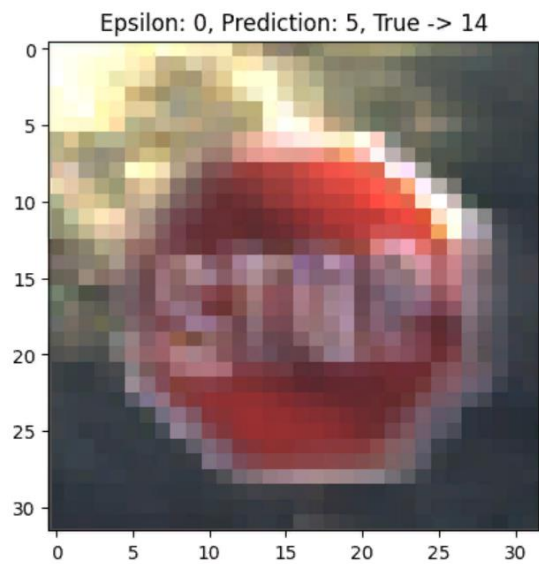
[6]

test = pd.read\_csv("Test.csv")
test\_imgs = test['Path'].values
data = []
y\_test = []
labels = test['ClassId'].values.tolist()
i = -1
for img in test\_imgs:
i += 1
if labels[i] != 14:
continue
img = image.load\_img(img, target\_size=(32, 32))
img\_array = image.img\_to\_array(img)
img\_array = img\_array / 255
data.append(img\_array)
y\_test.append(labels[i])
data = np.array(data)
y\_test = np.array(y\_test)
y\_test = to\_categorical(y\_test, 43)

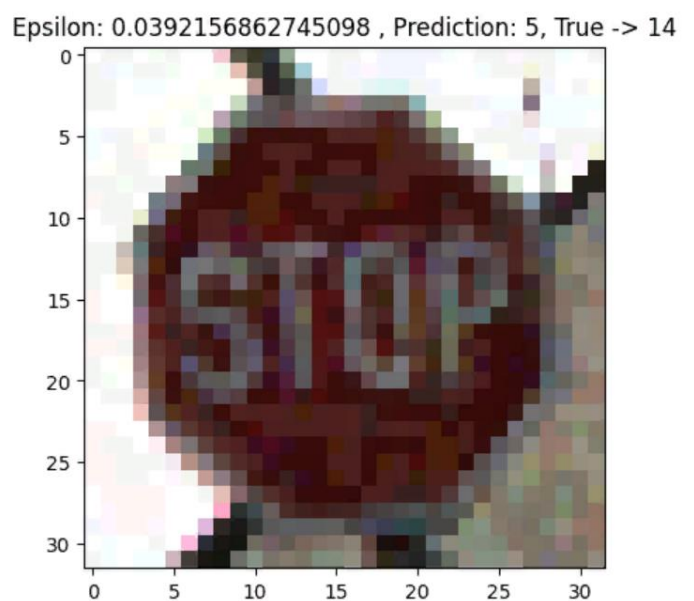

2)Построим 5 примеров и соответствующие атакующие FGSM примеры;








3)Повторим атаку PGD;







Epsilon: 0, Prediction: 5, True -> 14




Epsilon: 0.0392156862745098 , Prediction: 5, True -> 14




Epsilon: 0, Prediction: 5, True -> 14



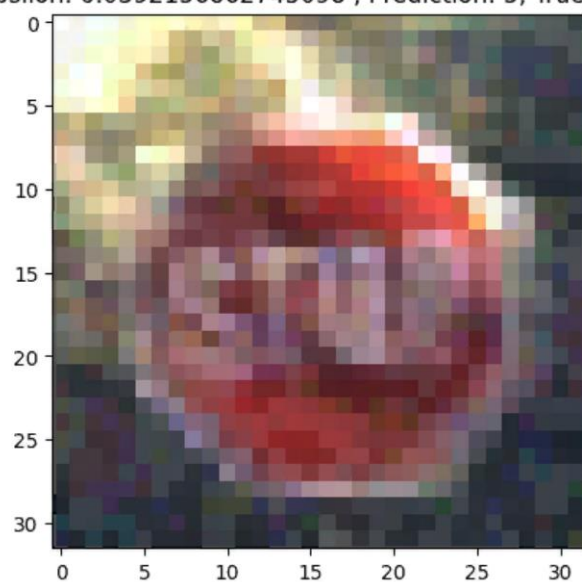
Epsilon: 0.0392156862745098 , Prediction: 5, True -> 14





Epsilon: 0, Prediction: 5, True -> 14



Epsilon: 0.0392156862745098 , Prediction: 5, True -> 14





4) Сохраним все значения точности идентично приведённым ниже и выведем значения в таблицу;

```

✓ 35
CEK. [14] attack_pgd = ProjectedGradientDescent(estimator=classifier, eps=1/255, max_iter=4, verbose=False)
adv_accuracises_pgd = []
x_test_adv = attack_pgd.generate(x_test, t_classes)
loss, accuracy = model.evaluate(x_test_adv, y_test)
adv_accuracises_pgd.append(accuracy)
resnetpgd_0=accuracy

attack_pgd = ProjectedGradientDescent(estimator=classifier, eps=3/255, max_iter=4, verbose=False)
adv_accuracises_pgd = []
x_test_adv = attack_pgd.generate(x_test, t_classes)
loss, accuracy = model.evaluate(x_test_adv, y_test)
adv_accuracises_pgd.append(accuracy)
resnetpgd_1=accuracy

attack_pgd = ProjectedGradientDescent(estimator=classifier, eps=5/255, max_iter=4, verbose=False)
adv_accuracises_pgd = []
x_test_adv = attack_pgd.generate(x_test, t_classes)
loss, accuracy = model.evaluate(x_test_adv, y_test)
adv_accuracises_pgd.append(accuracy)
resnetpgd_2=accuracy

attack_pgd = ProjectedGradientDescent(estimator=classifier, eps=10/255, max_iter=4, verbose=False)
adv_accuracises_pgd = []
x_test_adv = attack_pgd.generate(x_test, t_classes)
loss, accuracy = model.evaluate(x_test_adv, y_test)
adv_accuracises_pgd.append(accuracy)
resnetpgd_3=accuracy

attack_pgd = ProjectedGradientDescent(estimator=classifier, eps=20/255, max_iter=4, verbose=False)
adv_accuracises_pgd = []
x_test_adv = attack_pgd.generate(x_test, t_classes)
loss, accuracy = model.evaluate(x_test_adv, y_test)
adv_accuracises_pgd.append(accuracy)
resnetpgd_4=accuracy

attack_pgd = ProjectedGradientDescent(estimator=classifier, eps=50/255, max_iter=4, verbose=False)
adv_accuracises_pgd = []
x_test_adv = attack_pgd.generate(x_test, t_classes)
loss, accuracy = model.evaluate(x_test_adv, y_test)
adv_accuracises_pgd.append(accuracy)
resnetpgd_5=accuracy

attack_pgd = ProjectedGradientDescent(estimator=classifier, eps=80/255, max_iter=4, verbose=False)
adv_accuracises_pgd = []
x_test_adv = attack_pgd.generate(x_test, t_classes)
loss, accuracy = model.evaluate(x_test_adv, y_test)
adv_accuracises_pgd.append(accuracy)
resnetpgd_6=accuracy

```

```

from tabulate import tabulate

table = [{"Sign - Attack", "Adversarial images e=1/255", "Adversarial images e=3/255", "Adversarial images e=5/255", "Adversarial images e=10/255", "Adversarial images e=20/255", "Adversarial images e=50/255", "Adversarial images e=80/255"},
         ["Stop Sign - FGSM", resnetfgsm_0, resnetfgsm_1, resnetfgsm_2, resnetfgsm_3, resnetfgsm_4, resnetfgsm_5, resnetfgsm_6],
         ["Stop Sign - PGD", resnetpgd_0, resnetpgd_1, resnetpgd_2, resnetpgd_3, resnetpgd_4, resnetpgd_5, resnetpgd_6],]

table1 = tabulate(table, headers="firstrow", tablefmt="grid")
print(table1)

```

Sign - Attack	Adversarial images e=1/255	Adversarial images e=3/255	Adversarial images e=5/255	Adversarial images e=10/255	Adversarial images e=20/255	Adversarial images e=50/255	Adversarial images e=80/255
Stop Sign - FGSM	0.96	0.92	0.64	0.3	0.03	0	0
Stop Sign - PGD	0.97	0.93	0.86	0.77	0.45	0.1	0.04

Вывод:

В задание 2 по графикам можно определить, что для полученной модели, атака FGSM показала себя лучше, точность падает более медленно и

сбалансированно, что позволяет применять значения `epsilon`s точно и в большом диапазоне.

В задании 3 было получено обратное, метод FGSM показал себя хуже. В данном случае PGD была более стабильной и отлично справилась на целевой атаке. Наиболее оптимальное значение было получено при  $\epsilon = 20/255$ .