



МИНОБРНАУКИ РОССИИ  
Федеральное государственное бюджетное образовательное учреждение  
высшего образования  
**«МИРЭА – Российский технологический университет»**  
**РТУ МИРЭА**

---

ИКБ направление «Киберразведка и противодействие угрозам с применением технологий искусственного интеллекта» 10.04.01

Кафедра КБ-4 «Интеллектуальные системы информационной безопасности»

**Практическая работа №6**

по дисциплине

«Анализ защищённости систем искусственного интеллекта»

Группа:  
ББМО-01-22  
Выполнила:  
Огольцова Н.Д.

Проверил:  
Спирин А.А.

Москва 2023

## Ход работы:

1)Выполним загрузку необходимых библиотек;

```
✓ [1] import numpy as np
8 import matplotlib.pyplot as plt
%K import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import transforms, datasets
```

2)Задаём нормализующие преобразования и подгружаем датасет MNIST, разбиваем данные и выводим получившиеся значения;

```
✓ [2] transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.0,), (1.0,))])
%K

dataset = datasets.MNIST(root = './data', train=True, transform = transform, download=True)

train_set, val_set = torch.utils.data.random_split(dataset, [50000, 10000])
test_set = datasets.MNIST(root = './data', train=False, transform = transform, download=True)

train_loader = torch.utils.data.DataLoader(train_set, batch_size=1, shuffle=True)
val_loader = torch.utils.data.DataLoader(val_set, batch_size=1, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_set, batch_size=1, shuffle=True)

print("Training data:", len(train_loader), "Validation data:", len(val_loader), "Test data: ", len(test_loader))

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to ./data/MNIST/raw/train-images-idx3-ubyte.gz
100%|██████████| 9912422/9912422 [00:00<00:00, 120948579.17it/s]
Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to ./data/MNIST/raw/train-labels-idx1-ubyte.gz
100%|██████████| 28881/28881 [00:00<00:00, 113316832.39it/s]
Extracting ./data/MNIST/raw/train-labels-idx1-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw/t10k-images-idx3-ubyte.gz
100%|██████████| 1648877/1648877 [00:00<00:00, 22928621.76it/s]
Extracting ./data/MNIST/raw/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz
100%|██████████| 4542/4542 [00:00<00:00, 19990061.67it/s]
Extracting ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw

Training data: 50000 Validation data: 10000 Test data: 10000
```

3)Создаём класс НС на основе фреймворка torch и проверяем его работоспособность;

```

[4] class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, 3, 1)
        self.conv2 = nn.Conv2d(32, 64, 3, 1)
        self.dropout1 = nn.Dropout2d(0.25)
        self.dropout2 = nn.Dropout2d(0.5)
        self.fc1 = nn.Linear(9216, 128)
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = self.dropout1(x)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.dropout2(x)
        x = self.fc2(x)
        output = F.log_softmax(x, dim=1)
        return output

model = Net().to(device)

```

4) Создаём сначала оптимизатор, функцию потерь и потом трейнер сети;

```

[5] optimizer = optim.Adam(model.parameters(), lr=0.0001, betas=(0.9, 0.999))
    criterion = nn.NLLLoss()
    scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min', factor=0.1, patience=3)

```

5) Создадим функцию для обучения сети;

```

[6] def fit(model, device, train_loader, val_loader, epochs):
    data_loader = {'train': train_loader, 'val': val_loader}
    print("Fitting the model...")
    train_loss, val_loss = [], []
    for epoch in range(epochs):
        loss_per_epoch, val_loss_per_epoch = 0, 0
        for phase in ('train', 'val'):
            for i, data in enumerate(data_loader[phase]):
                input, label = data[0].to(device), data[1].to(device)
                output = model(input)
                #Вычисление выходных потерь
                loss = criterion(output, label)
                if phase == 'train':
                    optimizer.zero_grad()
                    #grad calc w.r.t Loss func
                    loss.backward()
                    #Обновление весов
                    optimizer.step()
                    loss_per_epoch += loss.item()
                else:
                    val_loss_per_epoch += loss.item()
            scheduler.step(val_loss_per_epoch/len(val_loader))
        print("Epoch: {} Loss: {} Val_Loss: {}".format(epoch+1, loss_per_epoch/len(train_loader), val_loss_per_epoch/len(val_loader)))
        train_loss.append(loss_per_epoch/len(train_loader))
        val_loss.append(val_loss_per_epoch/len(val_loader))
    return train_loss, val_loss

```

6) Обучим модель и выведем получившиеся значения;

✓  
27  
МИН.



```
loss, val_loss=fit(model,device,train_loader,val_loader,10)
```


Fitting the model...

/usr/local/lib/python3.10/dist-packages/torch/nn/functional.py:1345: UserWarning:   
warnings.warn(warn\_msg)

```
Epoch: 1 Loss: 0.2888687746945781 Val_Loss: 0.13673646724928057  
Epoch: 2 Loss: 0.10697086884837405 Val_Loss: 0.10052042624421069  
Epoch: 3 Loss: 0.08423662934597663 Val_Loss: 0.0830532077599337  
Epoch: 4 Loss: 0.0735847675861095 Val_Loss: 0.07899737361183146  
Epoch: 5 Loss: 0.06507188670661525 Val_Loss: 0.07766588417880915  
Epoch: 6 Loss: 0.061304444135245634 Val_Loss: 0.06547614247407182  
Epoch: 7 Loss: 0.056818889675883055 Val_Loss: 0.06705326536982777  
Epoch: 8 Loss: 0.05457629382752491 Val_Loss: 0.07752687670957667  
Epoch: 9 Loss: 0.053148277264218544 Val_Loss: 0.07033568611625667  
Epoch: 10 Loss: 0.051392144952551407 Val_Loss: 0.07953557262959682
```

7) Построим графики потерь при обучении и валидации в зависимости от эпохи;

```
fig = plt.figure(figsize=(5,5))  
plt.plot(np.arange(1,11), loss, "*-",label="Loss")  
plt.plot(np.arange(1,11), val_loss,"o-",label="Val Loss")  
plt.xlabel("Num of epochs")  
plt.legend()  
plt.show()
```



8) Создадим функции атак FGSM, I-FGSM, MI-FGSM;

```

def fgsm_attack(input,epsilon,data_grad):
    pert_out = input + epsilon*data_grad.sign()
    pert_out = torch.clamp(pert_out, 0, 1)
    return pert_out

def ifgsm_attack(input,epsilon,data_grad):
    iter = 10
    alpha = epsilon/iter
    pert_out = input
    for i in range(iter-1):
        pert_out = pert_out + alpha*data_grad.sign()
        pert_out = torch.clamp(pert_out, 0, 1)

    if torch.norm((pert_out-input),p=float('inf')) > epsilon:
        break
    return pert_out

def mifgsm_attack(input,epsilon,data_grad):
    iter=10
    decay_factor=1.0
    pert_out = input
    alpha = epsilon/iter
    g=0
    for i in range(iter-1):
        g = decay_factor*g + data_grad/torch.norm(data_grad,p=1)
        pert_out = pert_out + alpha*torch.sign(g)
        pert_out = torch.clamp(pert_out, 0, 1)

    if torch.norm((pert_out-input),p=float('inf')) > epsilon:
        break
    return pert_out

```

9)Создадим функцию тестирования/проверки;

```

[13] def test(model,device,test_loader,epsilon,attack):
    correct = 0
    adv_examples = []
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        data.requires_grad = True
        output = model(data)
        init_pred = output.max(1, keepdim=True)[1]
        if init_pred.item() != target.item():
            continue
        loss = F.nll_loss(output, target)
        model.zero_grad()
        loss.backward()
        data_grad = data.grad.data

        if attack == "fgsm":
            perturbed_data = fgsm_attack(data,epsilon,data_grad)
        elif attack == "ifgsm":
            perturbed_data = ifgsm_attack(data,epsilon,data_grad)
        elif attack == "mifgsm":
            perturbed_data = mifgsm_attack(data,epsilon,data_grad)

        output = model(perturbed_data)
        final_pred = output.max(1, keepdim=True)[1]
        if final_pred.item() == target.item():
            correct += 1
            if (epsilon == 0) and (len(adv_examples) < 5):
                adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
                adv_examples.append( (init_pred.item(), final_pred.item(), adv_ex) )
        else:
            if len(adv_examples) < 5:
                adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
                adv_examples.append( (init_pred.item(), final_pred.item(), adv_ex) )

    final_acc = correct/float(len(test_loader))
    print("Epsilon: {} \t Test Accuracy = {} / {} = {}".format(epsilon, correct, len(test_loader), final_acc))

    return final_acc, adv_examples

```

10) Построим графики ассурасу атак и выведем примеры выполненных атак в зависимости от значения epsilon;

```



epsilons = [0,0.007,0.01,0.02,0.03,0.05,0.1,0.2,0.3]
for attack in ("fgsm","ifgsm","mifgsm"):
    accuracies = []
    examples = []
    for eps in epsilons:
        acc, ex = test(model, device,test_loader,eps,attack)
        accuracies.append(acc)
        examples.append(ex)

    plt.figure(figsize=(5,5))
    plt.plot(epsilons, accuracies, "*-")
    plt.title(attack)
    plt.xlabel("Epsilon")
    plt.ylabel("Accuracy")
    plt.show()



    cnt = 0
    plt.figure(figsize=(8,10))
    for i in range(len(epsilons)):
        for j in range(len(examples[i])):
            cnt += 1
            plt.subplot(len(epsilons),len(examples[0]),cnt)
            plt.xticks([], [])
            plt.yticks([], [])
            if j == 0:
                plt.ylabel("Eps: {}".format(epsilons[i]), fontsize=14)
            orig,adv,ex = examples[i][j]
            plt.title("{} -> {}".format(orig, adv))
            plt.imshow(ex, cmap="gray")
    plt.tight_layout()
    plt.show()

```



Epsilon: 0      Test Accuracy = 9713 / 10000 = 0.9713  
 Epsilon: 0.007      Test Accuracy = 9685 / 10000 = 0.9685  
 Epsilon: 0.01      Test Accuracy = 9680 / 10000 = 0.968  
 Epsilon: 0.02      Test Accuracy = 9604 / 10000 = 0.9604  
 Epsilon: 0.03      Test Accuracy = 9547 / 10000 = 0.9547  
 Epsilon: 0.05      Test Accuracy = 9405 / 10000 = 0.9405  
 Epsilon: 0.1      Test Accuracy = 8735 / 10000 = 0.8735  
 Epsilon: 0.2      Test Accuracy = 6559 / 10000 = 0.6559  
 Epsilon: 0.3      Test Accuracy = 4063 / 10000 = 0.4063



Epsilon: 0      Test Accuracy = 9715 / 10000 = 0.9715  
 Epsilon: 0.007      Test Accuracy = 9693 / 10000 = 0.9693  
 Epsilon: 0.01      Test Accuracy = 9701 / 10000 = 0.9701  
 Epsilon: 0.02      Test Accuracy = 9667 / 10000 = 0.9667  
 Epsilon: 0.03      Test Accuracy = 9567 / 10000 = 0.9567  
 Epsilon: 0.05      Test Accuracy = 9458 / 10000 = 0.9458  
 Epsilon: 0.1      Test Accuracy = 8880 / 10000 = 0.888  
 Epsilon: 0.2      Test Accuracy = 7042 / 10000 = 0.7042  
 Epsilon: 0.3      Test Accuracy = 4804 / 10000 = 0.4804



Epsilon: 0 Test Accuracy = 9720 / 10000 = 0.972  
 Epsilon: 0.007 Test Accuracy = 9676 / 10000 = 0.9676  
 Epsilon: 0.01 Test Accuracy = 9676 / 10000 = 0.9676  
 Epsilon: 0.02 Test Accuracy = 9623 / 10000 = 0.9623  
 Epsilon: 0.03 Test Accuracy = 9580 / 10000 = 0.958  
 Epsilon: 0.05 Test Accuracy = 9435 / 10000 = 0.9435  
 Epsilon: 0.1 Test Accuracy = 8894 / 10000 = 0.8894  
 Epsilon: 0.2 Test Accuracy = 7031 / 10000 = 0.7031  
 Epsilon: 0.3 Test Accuracy = 4751 / 10000 = 0.4751



11)Создадим 2 класса НС;

✓  
0  
сек.

```
class NetF(nn.Module):
    def __init__(self):
        super(NetF, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, 3, 1)
        self.conv2 = nn.Conv2d(32, 64, 3, 1)
        self.dropout1 = nn.Dropout2d(0.25)
        self.dropout2 = nn.Dropout2d(0.5)
        self.fc1 = nn.Linear(9216, 128)
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = self.dropout1(x)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.dropout2(x)
        x = self.fc2(x)
        return x
```

```
class NetF1(nn.Module):
    def __init__(self):
        super(NetF1, self).__init__()
        self.conv1 = nn.Conv2d(1, 16, 3, 1)
        self.conv2 = nn.Conv2d(16, 32, 3, 1)
        self.dropout1 = nn.Dropout2d(0.25)
        self.dropout2 = nn.Dropout2d(0.5)
        self.fc1 = nn.Linear(4608, 64)
        self.fc2 = nn.Linear(64, 10)

    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = self.dropout1(x)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.dropout2(x)
        x = self.fc2(x)
        return x
```



## 12) Переопределим с учётом этого функцию обучения и тестирования;

```
def fit(model, device, optimizer, scheduler, criterion, train_loader, val_loader, Temp, epochs):
    data_loader = {'train': train_loader, 'val': val_loader}
    print("Fitting the model...")
    train_loss, val_loss = [], []
    for epoch in range(epochs):
        loss_per_epoch, val_loss_per_epoch = 0, 0
        for phase in ('train', 'val'):
            for i, data in enumerate(data_loader[phase]):
                input, label = data[0].to(device), data[1].to(device)
                output = model(input)
                output = F.log_softmax(output/Temp, dim=1)
                #расчет потерь на выходе
                loss = criterion(output, label)
                if phase == 'train':
                    optimizer.zero_grad()
                    #grad calc w.r.t Loss func
                    loss.backward()
                    #обновление весов
                    optimizer.step()
                    loss_per_epoch += loss.item()
                else:
                    val_loss_per_epoch += loss.item()
            scheduler.step(val_loss_per_epoch/len(val_loader))
        print("Epoch: {} Loss: {} Val_Loss: {}".format(epoch+1, loss_per_epoch/len(train_loader), val_loss_per_epoch/len(val_loader)))
        train_loss.append(loss_per_epoch/len(train_loader))
        val_loss.append(val_loss_per_epoch/len(val_loader))
    return train_loss, val_loss
```

```
def test(model, device, test_loader, epsilon, Temp, attack):
    correct = 0
    adv_examples = []
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        data.requires_grad = True
        output = model(data)
        output = F.log_softmax(output/Temp, dim=1)
        init_pred = output.max(1, keepdim=True)[1]
        if init_pred.item() != target.item():
            continue
        loss = F.nll_loss(output, target)
        model.zero_grad()
        loss.backward()
        data_grad = data.grad.data

        if attack == "fgsm":
            perturbed_data = fgsm_attack(data, epsilon, data_grad)
        elif attack == "ifgsm":
            perturbed_data = ifgsm_attack(data, epsilon, data_grad)
        elif attack == "mifgsm":
            perturbed_data = mifgsm_attack(data, epsilon, data_grad)

        output = model(perturbed_data)
        final_pred = output.max(1, keepdim=True)[1]
        if final_pred.item() == target.item():
            correct += 1
            if (epsilon == 0) and (len(adv_examples) < 5):
                adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
                adv_examples.append( (init_pred.item(), final_pred.item(), adv_ex) )
        else:
            if len(adv_examples) < 5:
                adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
                adv_examples.append( (init_pred.item(), final_pred.item(), adv_ex) )

    final_acc = correct/float(len(test_loader))
    print("Epsilon: {} \t Test Accuracy = {} / {} = {}".format(epsilon, correct, len(test_loader), final_acc))

    return final_acc, adv_examples
```

## 13) Создадим функцию защиты методом дистилляции;

```

def defense(device,train_loader,val_loader,test_loader,epochs,Temp,epsilons):

    modelF = NetF().to(device)
    optimizerF = optim.Adam(modelF.parameters(),lr=0.0001, betas=(0.9, 0.999))
    schedulerF = optim.lr_scheduler.ReduceLROnPlateau(optimizerF, mode='min', factor=0.1, patience=3)

    modelF1 = NetF1().to(device)
    optimizerF1 = optim.Adam(modelF1.parameters(),lr=0.0001, betas=(0.9, 0.999))
    schedulerF1 = optim.lr_scheduler.ReduceLROnPlateau(optimizerF1, mode='min', factor=0.1, patience=3)

    criterion = nn.NLLLoss()

    lossF,val_lossF=fit(modelF,device,optimizerF,schedulerF,criterion,train_loader,val_loader,Temp,epochs)
    fig = plt.figure(figsize=(5,5))
    plt.plot(np.arange(1,epochs+1), lossF, "--",label="Loss")
    plt.plot(np.arange(1,epochs+1), val_lossF,"o-",label="Val Loss")
    plt.title("Network F")
    plt.xlabel("Num of epochs")
    plt.legend()
    plt.show()

    #Преобразование целевых меток в гибкие метки
    for data in train_loader:
        input, label = data[0].to(device),data[1].to(device)
        softlabel = F.log_softmax(modelF(input),dim=1)
        data[1] = softlabel

```

```

lossF1,val_lossF1=fit(modelF1,device,optimizerF1,schedulerF1,criterion,train_loader,val_loader,Temp,epochs)
fig = plt.figure(figsize=(5,5))
plt.plot(np.arange(1,epochs+1), lossF1, "--",label="Loss")
plt.plot(np.arange(1,epochs+1), val_lossF1,"o-",label="Val Loss")
plt.title("Network F'")
plt.xlabel("Num of epochs")
plt.legend()
plt.show()

model = NetF1().to(device)
model.load_state_dict(modelF1.state_dict())
for attack in ("fgsm","ifgsm","mifgsm"):
    accuracies = []
    examples = []
    for eps in epsilons:
        acc, ex = test(model,device,test_loader,eps,1,"fgsm")
        accuracies.append(acc)
        examples.append(ex)

    plt.figure(figsize=(5,5))
    plt.plot(epsilons, accuracies, "--")
    plt.title(attack)
    plt.xlabel("Epsilon")
    plt.ylabel("Accuracy")
    plt.show()

    cnt = 0
    plt.figure(figsize=(8,10))
    for i in range(len(epsilons)):
        for j in range(len(examples[i])):
            cnt += 1
            plt.subplot(len(epsilons),len(examples[0]),cnt)
            plt.xticks([], [])
            plt.yticks([], [])
            if j == 0:
                plt.ylabel("Eps: {}".format(epsilons[i]), fontsize=14)
            orig,adv,ex = examples[i][j]
            plt.title("{} -> {}".format(orig, adv))
            plt.imshow(ex, cmap="gray")
    plt.tight_layout()
    plt.show()

```

14)Получим результаты оценки уже защищенных сетей.

```

Temp=100
epochs=10
epsilons=[0,0.007,0.01,0.02,0.03,0.05,0.1,0.2,0.3]
defense(device,train_loader,val_loader,test_loader,epochs,Temp,epsilons)



```

Fitting the model...



Epoch: 1 Loss: 0.5183894712342154 Val\_Loss: 0.35013818526350576  
Epoch: 2 Loss: 0.3003614776288117 Val\_Loss: 0.252035639292655  
Epoch: 3 Loss: 0.21513449889052666 Val\_Loss: 0.19123540499878997  
Epoch: 4 Loss: 0.16343666033960325 Val\_Loss: 0.14716214833448865  
Epoch: 5 Loss: 0.13425891620466573 Val\_Loss: 0.12984373547466652  
Epoch: 6 Loss: 0.11724456572161095 Val\_Loss: 0.11608159823280148  
Epoch: 7 Loss: 0.10449432968694738 Val\_Loss: 0.09994257388982973  
Epoch: 8 Loss: 0.0953535752760143 Val\_Loss: 0.10708337319088777  
Epoch: 9 Loss: 0.08788794863414845 Val\_Loss: 0.09572349952923845  
Epoch: 10 Loss: 0.08311829223801133 Val\_Loss: 0.08903677334889062

Fitting the model...



Epoch: 1 Loss: 0.6932818498760024 Val\_Loss: 0.49934649001323944  
Epoch: 2 Loss: 0.46901066487247844 Val\_Loss: 0.4374905162987688  
Epoch: 3 Loss: 0.422528915409889 Val\_Loss: 0.41707982867373905  
Epoch: 4 Loss: 0.37601292758400523 Val\_Loss: 0.3597610071243918  
Epoch: 5 Loss: 0.3387371742089258 Val\_Loss: 0.32266275289950297  
Epoch: 6 Loss: 0.2994880652114021 Val\_Loss: 0.28805370525927454  
Epoch: 7 Loss: 0.26132685050939375 Val\_Loss: 0.2564337362716436  
Epoch: 8 Loss: 0.22919203091141782 Val\_Loss: 0.22150387843409927  
Epoch: 9 Loss: 0.20354706166494133 Val\_Loss: 0.19780019457405493  
Epoch: 10 Loss: 0.18127455468622433 Val\_Loss: 0.1758993640845126





Epsilon: 0 Test Accuracy = 9251 / 10000 = 0.9251  
Epsilon: 0.007 Test Accuracy = 9246 / 10000 = 0.9246  
Epsilon: 0.01 Test Accuracy = 9285 / 10000 = 0.9285  
Epsilon: 0.02 Test Accuracy = 9274 / 10000 = 0.9274  
Epsilon: 0.03 Test Accuracy = 9280 / 10000 = 0.928  
Epsilon: 0.05 Test Accuracy = 9227 / 10000 = 0.9227  
Epsilon: 0.1 Test Accuracy = 9197 / 10000 = 0.9197  
Epsilon: 0.2 Test Accuracy = 9139 / 10000 = 0.9139  
Epsilon: 0.3 Test Accuracy = 9097 / 10000 = 0.9097



Epsilon: 0 Test Accuracy = 9281 / 10000 = 0.9281  
 Epsilon: 0.007 Test Accuracy = 9276 / 10000 = 0.9276  
 Epsilon: 0.01 Test Accuracy = 9245 / 10000 = 0.9245  
 Epsilon: 0.02 Test Accuracy = 9253 / 10000 = 0.9253  
 Epsilon: 0.03 Test Accuracy = 9214 / 10000 = 0.9214  
 Epsilon: 0.05 Test Accuracy = 9267 / 10000 = 0.9267  
 Epsilon: 0.1 Test Accuracy = 9179 / 10000 = 0.9179  
 Epsilon: 0.2 Test Accuracy = 9116 / 10000 = 0.9116  
 Epsilon: 0.3 Test Accuracy = 9085 / 10000 = 0.9085



Epsilon: 0 Test Accuracy = 9209 / 10000 = 0.9209  
 Epsilon: 0.007 Test Accuracy = 9256 / 10000 = 0.9256  
 Epsilon: 0.01 Test Accuracy = 9277 / 10000 = 0.9277  
 Epsilon: 0.02 Test Accuracy = 9246 / 10000 = 0.9246  
 Epsilon: 0.03 Test Accuracy = 9277 / 10000 = 0.9277  
 Epsilon: 0.05 Test Accuracy = 9245 / 10000 = 0.9245  
 Epsilon: 0.1 Test Accuracy = 9189 / 10000 = 0.9189  
 Epsilon: 0.2 Test Accuracy = 9151 / 10000 = 0.9151  
 Epsilon: 0.3 Test Accuracy = 9080 / 10000 = 0.908



Вывод:

При увеличении значения переменной `epsilon` у атак - точность падает практически вдвое, падает от значения 0.97 в среднем до 0.45.

Тогда как при атаке на защищённые сети значение остаётся в пределах от 0.92 - 0.90.

Также можно заметить, что у защищённых сетей (при значении `epsilon = 0` -> атака отсутствует) значение точности снизилось.